

## 第 0019 讲 6 内存管理 meminfo&zoneinfo 信息分析

### 内存管理 meminfo&zoneinfo 信息分析

#### 一、详细分析/proc/meminfo 节点信息

```
vico@ubuntu: ~  
File Edit View Search Terminal Help  
vico@ubuntu:~$ cat /proc/meminfo  
MemTotal:      16362520 kB  
MemFree:       14309592 kB  
MemAvailable:  14846336 kB  
Buffers:       111796 kB  
Cached:        617944 kB  
SwapCached:    0 kB  
Active:        1152444 kB  
Inactive:      389552 kB  
Active(anon):  805940 kB  
Inactive(anon): 6416 kB  
Active(file):  346504 kB  
Inactive(file): 383136 kB  
Unevictable:   16 kB  
Mlocked:       16 kB  
SwapTotal:     2097148 kB  
SwapFree:      2097148 kB  
Dirty:         196 kB  
Writeback:     0 kB  
AnonPages:     812284 kB  
Mapped:        236452 kB  
Shmem:         7772 kB  
KReclaimable:  121192 kB  
Slab:          196964 kB
```

具体系统调用实现，在 Linux 内核源码函数设计如下：

```
fs > proc > C meminfo.c > meminfo_proc_show(seq_file *, void *)  
34 static int meminfo_proc_show(struct seq_file *m, void *v)  
35 {  
36     struct sysinfo i;  
37     unsigned long committed;  
38     long cached;  
39     long available;  
40     unsigned long pages[NR_LRU_LISTS];  
41     unsigned long sreclaimable, sunreclaim;  
42     int lru;  
43  
44     si_meminfo(&i);  
45     si_swapinfo(&i);  
46     committed = percpu_counter_read_positive(&vm_committed_as);  
47  
48     cached = global_node_page_state(NR_FILE_PAGES) -  
49             total_swapcache_pages() - i.bufferram;  
50     if (cached < 0)  
51         cached = 0;  
52  
53     for (lru = LRU_BASE; lru < NR_LRU_LISTS; lru++)  
54         pages[lru] = global_node_page_state(NR_LRU_BASE + lru);  
55  
56     available = si_mem_available();  
57     sreclaimable = global_node_page_state(NR_SLAB_RECLAIMABLE);  
58     sunreclaim = global_node_page_state(NR_SLAB_UNRECLAIMABLE);  
59  
60     show_val_kb(m, "MemTotal:", i.totalram);  
61     show_val_kb(m, "MemFree:", i.freeram);
```

```
root@ubuntu:/home/vico# cat /proc/meminfo
```

MemTotal: 16362520 kB // 表示当前系统可用物理内存总数

MemFree: 14455324 kB // 表示当前系统剩余空闲物理内存

// 表示系统中可使用页面的数量（空闲页面、文件映射页面、可回收的页面）

公式：

可使用页面数量=空闲页面+文件映射页面+可回收页面-系统保留的页面

MemAvailable: 14873756 kB

Buffers: 54688 kB // 用于块层的缓存

Cached: 585220 kB // 用于页面高速缓存的页面

SwapCached: 0 kB // 统计交换缓存的数量

// 活跃的匿名页面（LRU\_ACTIVE\_ANON）和活跃的文件映射页面（LRU\_ACTIVE\_FILE）

Active: 1091584 kB

// 不活跃的匿名页面（LRU\_INACTIVE\_ANON）和不活跃的文件映射页面

（LRU\_INACTIVE\_FILE）

Inactive: 372708 kB

Active(anon): 819572 kB // 活跃的匿名页面

Inactive(anon): 6376 kB // 不活跃的匿名碳

Active(file): 272012 kB // 活跃的文件映射页面

Inactive(file): 366332 kB // 不活跃的文件映射碳

Unevictable: 16 kB // 不能回收的碳

Mlocked: 16 kB // 不会被交换到交换分区的碳

SwapTotal: 2097148 kB // 交换分析的大小

SwapFree: 2097148 kB // 交换分区的空闲空间大小

Dirty: 776 kB // 脏页的数量

Writeback: 0 kB // 正在回写的页面数量

AnonPages: 824428 kB // 统计有反向映射的页面

Mapped: 238420 kB // 统计所有映射到用户地址空间的内存缓存页面

Shmem: 7724 kB // 共享内存页面的数量

KReclaimable: 67156 kB // Linux 内核可回收的内存, 包含回收的 slab 页面

Slab: 143680 kB // 所有 slab 页面

SReclaimable: 67156 kB // 可回收的 slab 页面

SUnreclaim: 76524 kB // 不可回收的 slab 页面

KernelStack: 12480 kB // 所有进程内核栈大小

PageTables: 42316 kB // 所有用于页表的页面数量

NFS\_Unstable: 0 kB // 在 NFS 中, 发送到服务器端但是还没有写入磁盘的页面

Bounce: 0 kB // 退回

WritebackTmp: 0 kB // 回写过程中的使用临时缓存

CommitLimit: 10278408 kB // 虚拟内存限值

Committed\_AS: 4527116 kB // 表示系统已经分配的内存情况

VmallocTotal: 34359738367 kB // vmalloc 区域的总大小

VmallocUsed: 27596 kB // 已使用的 vmalloc 区域总大小

VmallocChunk: 0 kB // 最大的连续未被使用 vmalloc 区域

Percpu: 46080 kB // percpu 机制使用的页面, 由 percpu\_nr\_pages() 函数来统计

HardwareCorrupted: 0 kB // 当系统检测到内存的硬件故障时, 会把有问题的页面删除

AnonHugePages: 0 kB // 统计透明巨型页的数量

ShmemHugePages: 0 kB

ShmemPmdMapped: 0 kB // 使用透明巨型页并且映射到用户空间

FileHugePages: 0 kB

FilePmdMapped: 0 kB

CmaTotal: 0 kB

CmaFree: 0 kB

HugePages\_Total: 0

HugePages\_Free: 0

HugePages\_Rsvd: 0

HugePages\_Surp: 0

Hugepagesize: 2048 kB

Hugetlb: 0 kB

DirectMap4k: 161600 kB

DirectMap2M: 5081088 kB

DirectMap1G: 13631488 kB

root@ubuntu:/home/vico#

## 二、详细分析/proc/zoneinfo 节点信息

```
root@ubuntu: /home/vico
File Edit View Search Terminal Help
root@ubuntu:/home/vico# cat /proc/zoneinfo
Node 0, zone DMA
per-node stats
  nr_inactive_anon 2161
  nr_active_anon 258983
  nr_inactive_file 94786
  nr_active_file 130957
  nr_unevictable 4
  nr_slab_reclaimable 31311
  nr_slab_unreclaimable 19771
  nr_isolated_anon 0
  nr_isolated_file 0
  workingset_nodes 0
  workingset_refault 0
  workingset_activate 0
  workingset_restore 0
  workingset_nodereclaim 0
  nr_anon_pages 260563
  nr_mapped 92841
  nr_file_pages 226329
  nr_dirty 176
  nr_writeback 0
  nr_writeback_temp 0
```

具体系统调用实现，在 Linux 内核源码函数设计如下：

```
mm > C vmstat.c > zoneinfo_show_print(seq_file *, pg_data_t *, zone *)
1544 static void zoneinfo_show_print(struct seq_file *m, pg_data_t *pgdat,
1545                                struct zone *zone)
1546 {
1547     int i;
1548     seq_printf(m, "Node %d, zone %8s", pgdat->node_id, zone->name);
1549     if (is_zone_first_populated(pgdat, zone)) {
1550         seq_printf(m, "\n per-node stats");
1551         for (i = 0; i < NR_VM_NODE_STAT_ITEMS; i++) {
1552             seq_printf(m, "\n %-12s %lu",
1553                 vmstat_text[i + NR_VM_ZONE_STAT_ITEMS +
1554                     NR_VM_NUMA_STAT_ITEMS],
1555                 node_page_state(pgdat, i));
1556         }
1557     }
1558     seq_printf(m,
1559         "\n pages free %lu"
1560         "\n min %lu"
1561         "\n low %lu"
1562         "\n high %lu"
1563         "\n spanned %lu"
1564         "\n present %lu"
1565         "\n managed %lu",
```

```
root@ubuntu:/home/vico# cat /proc/zoneinfo
```

```
Node 0, zone      DMA
```

```
per-node stats
```

```
nr_inactive_anon 1592
```

```
nr_active_anon 206450
```

```
nr_inactive_file 96845
```

```
nr_active_file 84139
```

```
nr_unevictable 4
```

```
nr_slab_reclaimable 30210
```

```
nr_slab_unreclaimable 19212
```

```
nr_isolated_anon 0
```

```
nr_isolated_file 0
```

```
workingset_nodes 0
```

```
workingset_refault 0
```

```
workingset_activate 0
```

```
workingset_restore 0
```

```
workingset_nodereclaim 0
```

```
nr_anon_pages 207679
```

```
nr_mapped      59666
```

```
nr_file_pages 181354
```

```
nr_dirty        115
```

```
nr_writeback 0
```

```
nr_writeback_temp 0
```

```
nr_shmem        1930
```

```
nr_shmem_hugepages 0
```

```
nr_shmem_pmdmapped 0
```

```
nr_file_hugepages 0
```

```
nr_file_pmdmapped 0
```

```
nr_anon_transparent_hugepages 0
```

```
nr_unstable     0
```

```
nr_vmscan_write 0
```

```
nr_vmscan_immediate_reclaim 0
```

```
nr_dirtied      27385
```

```
nr_written       8105
```

```
nr_kernel_misc_reclaimable 0
```

```
pages free      3976
```

```
min             16
```

```
low             20
```

```
high           24
```

```
spanned        4095
```

```
present        3997
```

```
managed        3976
```

```
protection: (0, 2950, 15914, 15914, 15914)
```

```
nr_free_pages 3976
nr_zone_inactive_anon 0
nr_zone_active_anon 0
nr_zone_inactive_file 0
nr_zone_active_file 0
nr_zone_unevictable 0
nr_zone_write_pending 0
nr_mlock      0
nr_page_table_pages 0
nr_kernel_stack 0
nr_bounce     0
nr_zspages    0
nr_free_cma   0
numa_hit      0
numa_miss     0
numa_foreign  0
numa_interleave 0
numa_local    0
numa_other    0
pagesets
cpu: 0
        count: 0
        high:  0
        batch: 1
vm stats threshold: 2
node_unreclaimable:  0
start_pfn:           1
Node 0, zone DMA32
pages free           765443
        min           3128
        low           3910
        high          4692
        spanned      1044480
        present      782288
        managed      765904
        protection: (0, 0, 12964, 12964, 12964)
nr_free_pages 765443
nr_zone_inactive_anon 0
nr_zone_active_anon 0
nr_zone_inactive_file 0
nr_zone_active_file 0
nr_zone_unevictable 0
nr_zone_write_pending 0
nr_mlock      0
```



```

nr_page_table_pages 0
nr_kernel_stack 0
nr_bounce    0
nr_zspages   0
nr_free_cma   0
numa_hit     16
numa_miss    0
numa_foreign 0
numa_interleave 1
numa_local   16
numa_other    0
Pagesets // // 表示每个 CPU 内存分配器中每个 CPU 缓存的页面数据信息
cpu: 0
    count: 400
    high: 378
    batch: 63
vm stats threshold: 12
node_unreclaimable: 0 // 表示页面回收失败的次数
start_pfn: 4096 // 表示内存管理区的起始页帧号
Node 0, zone Normal
pages free 2806651
    min 13750
    low 17187
    high 20624
spanned 3407872
present 3407872
managed 3320750
protection: (0, 0, 0, 0, 0)
nr_free_pages 2806651
nr_zone_inactive_anon 1592
nr_zone_active_anon 206450
nr_zone_inactive_file 96845
nr_zone_active_file 84139
nr_zone_unevictable 4
nr_zone_write_pending 115
nr_mlock 4
nr_page_table_pages 10576
nr_kernel_stack 11904
nr_bounce 0
nr_zspages 0
nr_free_cma 0
numa_hit 877306
numa_miss 0
numa_foreign 0

```



```
numa_interleave 32733
numa_local      877306
numa_other      0
pagesets
cpu: 0
        count: 191
        high:  378
        batch: 63
vm stats threshold: 16
node_unreclaimable: 0
start_pfn:      1048576
Node 0, zone    Movable
pages free      0
        min      0
        low      0
        high     0
        spanned  0
        present  0
        managed  0
        protection: (0, 0, 0, 0, 0)
Node 0, zone    Device
pages free      0 // 表示此内存管理区中空闲页面的数量
        min      0 // 表示此内存管理区处于最低水位的页面数量
        low      0 // 表示此内存管理区处于低水位的页面数量
        high     0 // 表示此内存管理区处于高水位的页面数量
        spanned  0 // 表示此内存管理区包含的页面数量
        present  0 // 表示此内存管理区里实际管理的页面数量
        managed  0 // 表示此内存管理区中被伙伴系统管理的页面数量
        protection: (0, 0, 0, 0, 0) // 表示此内存管理区预留的内存
root@ubuntu:/home/vico#
```